

Diffusion Drafters for Speculative² Decoding: A Systems Study

Cashel Fitzgerald
Cornell University
cf566@cornell.edu

Tishya Khanna
Cornell University
tk732@cornell.edu

Abstract

Speculative decoding accelerates autoregressive language model inference by letting a cheap draft model propose multiple future tokens that a target model verifies in parallel. Speculative² Decoding (S²D) goes further: while the target verifies one draft, a separate draft device precomputes continuations for likely verification outcomes. This naturally invites a diffusion drafter, since a block-diffusion model can in principle generate an entire draft block in parallel. We study that idea in two forms. First, we integrate conventional masked and block-diffusion language models into the S²D draft interface. This negative result is the main lesson: practical block diffusion still requires many refinement forwards, often approaching the number of tokens in the block, so the drafter becomes the bottleneck. Our best traditional block-diffusion runs reached only about 20 tokens/s. Second, we prototype a DFlash-style hybrid that uses target activations from the verifier to produce a first draft block on S²D cache misses, while the normal autoregressive S²D drafter serves cache hits and fills the future speculation cache. This is architecturally more promising because DFlash drafts a block in a single conditioned pass, but our initial small-model prototype remains limited by systems overhead: on Qwen3-8B with a Qwen3-0.6B S²D drafter and Qwen3-8B-DFlash-b16, we measured 11.68 tokens/s, average accepted length 3.30, cache hit rate 0.60, DFlash seed latency 34.48 ms, and SSD tree-decode latency about 79 ms. Moving the prototype to a two-H100 Modal deployment and shortening only the S²D lookahead from 15 to 7 improved a matched diagnostic run from 31.15 to 32.82 tokens/s by reducing draft-side tree work, while still using the native 15-token DFlash seed. We conclude that diffusion drafters are only useful for S²D when draft latency is far below target verification latency, and that representative evaluation requires a larger verifier and less resource-constrained hardware.

1. Introduction

Autoregressive (AR) large language model inference is sequential: each generated token depends on the full prefix, and each step therefore requires another target-model forward pass. This is especially costly for verifier models in the tens of billions of parameters, where latency is dominated by memory movement and KV-cache traffic rather than by abundant GPU arithmetic. Speculative decoding (SD) reduces the number of serial target forwards by using a cheaper draft model to propose k future tokens and verifying those tokens in one target pass [2, 6]. The resulting accept/reject rule is lossless: the final sequence is distributed exactly as if it had been sampled from the target.

Speculative² Decoding (S²D), implemented as Saguario, attacks the next source of idle time [5]. In ordinary SD, the draft model waits while the target verifies, and the target then waits while the draft prepares the next block. S²D places the draft model on separate hardware and asks it to predict likely verification outcomes while the target is busy. If the actual outcome is one of the precomputed branches, the next draft is ready immediately. The method is attractive because it can improve the throughput-latency frontier without changing the verifier or giving up exact sampling.

This report asks whether diffusion language models are a good replacement for the AR drafter inside S²D. The motivation is straightforward. Discrete diffusion language models denoise masked tokens in parallel [1, 7], so a diffusion drafter appears to remove the inner AR bottleneck from the draft side. If one block forward could replace k serial AR draft steps, S²D could use its separate draft GPU more effectively and cache likely future branches faster.

Our experiments show that this idea is more subtle. Conventional masked or block-diffusion drafters do not behave like a one-pass block oracle. To get usable draft quality, they require several denoising/refinement forwards and, in the quality regime, the number of refinement passes can approach the number of tokens generated in the block. In S²D this cost is paid on the draft critical path and can dominate the very latency S²D is meant to hide. In our implementa-

tion, the traditional block-diffusion backend peaked around 20 tokens/s and was not competitive with AR draft baselines.

DFlash changes the trade-off [3]. Rather than asking a standalone diffusion model to infer the next block from tokens alone, DFlash conditions a lightweight block-diffusion drafter on hidden activations extracted from the target verifier. This target-conditioned design produces a whole speculative block in a single draft pass and has published speedups on Qwen3-8B and Qwen3-Coder-30B-A3B using SGLang [9]. We therefore extended our study beyond conventional block diffusion and implemented a DFlash-seeded S²D hybrid. We also decouple the DFlash block length from the S²D tree depth: the DFlash checkpoint still produces its native 15-token block, but the SSD verifier and tree cache may consume a shorter prefix of that block to reduce draft-side work.

Contributions. First, we integrated MDLM-style and BD3LM-style drafters into an S²D inference engine and measured why they do not provide useful speedups in this setting. Second, we implemented a DFlash hybrid in which DFlash serves S²D speculation-cache misses using fresh verifier activations, while the normal AR S²D drafter serves cache hits and prepares the future tree cache. Third, we ablate this design with a *universal drafter* variant that also runs DFlash on cache hits, using the cached continuation as a non-mask prior; this lets us isolate whether re-drafting hits adds signal or simply costs latency. Fourth, we identify the main systems bottleneck: diffusion drafting is useful only when the verifier is sufficiently slower than the draft path to amortize cross-GPU communication, scheduling, tree-cache construction, and fallback logic. On small verifier models, those overheads erase the advantage. Finally, we outline the larger-model experiment needed to decide whether DFlash+S²D is worthwhile under representative serving conditions.

2. Method

2.1. Background

Speculative decoding. Given prefix $x_{\leq t}$, a draft model M_d proposes $\hat{x}_{t+1:t+k}$ autoregressively. The target model M_t then evaluates the drafted positions in one forward pass and accepts each token according to the standard rejection-sampling rule $\min(1, p_{M_t}(\hat{x}_i)/p_{M_d}(\hat{x}_i))$ [2, 6]. At the first rejection, SD samples a corrected token from the residual distribution. Under this rule, SD changes latency but not the output distribution.

Speculative² decoding. S²D keeps the target and draft on separate devices [5]. While the target verifies the current speculation, the draft predicts likely verification outcomes (a, r) , where a is the number of accepted draft tokens and r

is the recovery or bonus token. It then precomputes continuations for those outcomes and stores them in a speculation cache. A cache hit eliminates the next drafting wait; a cache miss falls back to a synchronous or just-in-time draft.

Discrete diffusion drafters. Masked diffusion language models iteratively replace mask tokens with predictions [7]. Block diffusion adds a block-causal attention pattern, often called staircase attention, so that a sequence is generated block by block while positions inside the current block are denoised in parallel [1]. These models are bidirectional inside the open block, so only fully committed prefix blocks are safe to reuse in a KV cache; the partially filled open block must be recomputed with accepted tokens clamped and unresolved positions masked.

2.2. Traditional block diffusion in S²D

Our first integration adds a block draft backend with three samplers: mask-predict [4], remask, and first-hitting. Given a prefix and lookahead k , the backend constructs

$$[x_{\leq t}, \underbrace{\text{MASK}, \dots, \text{MASK}}_k]$$

and runs R refinement forwards. Each forward predicts logits for all masked positions, commits a subset of tokens according to the sampler, and reuses or recomputes the open block as required by the attention pattern. We implemented both full bidirectional attention and staircase attention with prefix-cache reuse for committed blocks.

This is the most direct diffusion replacement for an AR drafter, but it is not the right systems trade-off. In practice, useful drafts require multiple diffusion forwards. The draft cost is therefore not one block pass; it is approximately

$$T_{\text{draft}} \approx R \cdot T_{\text{diffusion_forward}},$$

plus cache construction and transfer overhead. When R becomes large enough to produce acceptable tokens, this can approach the cost of generating the block token-by-token with a small AR drafter. Since S²D only helps if the draft path can be hidden under target verification, this iterative diffusion cost becomes a bottleneck.

2.3. DFlash-seeded S²D

DFlash addresses the quality-latency problem by conditioning a lightweight block-diffusion drafter on target hidden states [3]. The target verifier already computes rich activations during prefill and verify; DFlash concatenates selected target-layer residual streams, injects those features into the draft model’s KV cache, and predicts the next block from a recovery token followed by mask tokens. The DFlash checkpoint we target, z-lab/Qwen3-8B-DFlash-b16, has block size 16, so our v1 path uses 15 native DFlash draft

tokens. For synchronous DFlash, this fixes the speculation lookahead at $k = 15$. For the DFlash+S²D hybrid, however, we separate the DFlash lookahead from the SSD lookahead. Cache misses always ask DFlash for the full 15-token block required by the checkpoint, while the SSD response may expose only the first k_{SSD} tokens and logits to the verifier. This keeps the DFlash checkpoint geometry unchanged while allowing the SSD tree depth, fan-out buffers, and tree-cache work to shrink.

We integrate DFlash with S²D as a miss handler rather than replacing the whole S²D loop. The hybrid works as follows.

1. The verifier runs the target model and returns logits, recovery tokens, and DFlash target activations.
2. The async draft process receives a speculation-cache request keyed by the sequence id, accepted depth, and recovery token.
3. On a cache hit, S²D returns the normal cached AR-draft branch.
4. On a cache miss, DFlash receives the verifier activation backlog, builds [recovery, MASK¹⁵], and produces a 15-token draft block using the target embedding and LM head. In the split-lookahead variant, the async response slices this block to the first k_{SSD} positions before target verification.
5. After the exact verifier accepts or rejects the proposed block, the scheduler appends only the verified target activations to the DFlash backlog and crops the DFlash cache back to the verified prefix.
6. In parallel, the existing S²D AR drafter continues constructing the tree cache for future likely verification outcomes.

The important design choice is that DFlash is used only when the S²D tree cache misses. Hits keep the original S²D fast path, where precomputed AR branches are returned immediately. Misses avoid the weakest fallback of ordinary S²D, namely waiting for a full just-in-time AR draft, and instead use the verifier-conditioned DFlash model to provide the next block.

2.4. Universal-drafter ablation

The baseline hybrid (V1) only runs DFlash on cache misses; hits return the cached AR branch verbatim. As an ablation we built a *universal drafter* variant (V2) that also invokes DFlash on hits, seeding its block as [recovery, $c_1, \dots, c_{15}$] with the cached tokens c_i instead of the V1 form [recovery, MASK, ...]. The same target-conditioned DFlash forward is used; on a miss V2 is identical

to V1. V2 additionally advances the DFlash per-sequence KV cache and activation backlog on hits (which V1 leaves untouched).

2.5. Correctness and v1 restrictions

Both block diffusion and DFlash are still draft mechanisms: the target verifier remains exact. At greedy temperature zero, matching AR outputs token-for-token is the correctness test. For stochastic decoding, the same SD residual rule would preserve the target distribution, but our implementation currently focuses on greedy decoding to isolate systems behavior.

The current DFlash+S²D prototype intentionally supports only batch size 1, two GPUs, Qwen3-style target activations, and greedy decoding. The verifier and drafter are placed on separate visible GPUs. The synchronous DFlash diagnostic path is single-GPU and keeps the DFlash-native $k = 15$ lookahead. The async hybrid additionally maintains DFlash KV state and target activation backlogs per sequence id, so multiple independent benchmark requests do not share stale DFlash cache state. This is sufficient to test the central systems question: can a one-pass target-conditioned diffusion block drafter make S²D cache misses cheaper than AR fallback?

3. Experiments

3.1. Setup

Reference SSD evaluation. The original S²D paper evaluates batch-size-one greedy decoding with the target on 4×H100 and, for SSD, the asynchronous draft on a separate H100 [5]. Its Llama setup uses Llama-3.1-70B-Instruct with a Llama-3.2-1B draft model; the Qwen appendix uses Qwen3-32B with a Qwen3-0.6B draft model [8]. The benchmark pool contains 128 prompts from each of HumanEval, Alpaca, GSM8K, and UltraFeedback, for 512 prompts total, with 512 output tokens per prompt.

Our local evaluation. Our diffusion experiments were constrained by available compute. We used shared A6000-class machines for implementation and smoke testing, and several larger runs were blocked by other processes occupying most GPU memory. The reported prototype numbers should therefore be read as diagnostic measurements, not as final throughput claims. This limitation matters: S²D is a latency-hiding method, so evaluating it on a small verifier or on busy GPUs can make communication and scheduler overhead look much worse than it would with a larger target. For follow-up diagnostics, we moved the Qwen3-8B DFlash+S²D experiments to a Modal deployment with persistent model and compiler caches. That app provides one-H100 entryptoints for AR and synchronous DFlash checks, two-H100 entryptoints for the async DFlash+S²D path, and

Table 1. Published S²D results from Kumar et al. [5]. Values are end-to-end decode throughput in tokens/s averaged over four datasets.

Model pair	AR	SD	SSD	SSD/AR
Llama 70B / 1B	54.7	161.8	255.8	4.68×
Qwen3 32B / 0.6B	88.8	136.8	203.8	2.29×

five-H100 entrypoints reserved for the original 70B SSD configuration.

Models. For the traditional block-diffusion path, we used the small Qwen-family stack available in the repository and the block/MDLM-style draft checkpoints supported by our backend. For the DFlash hybrid, we used Qwen3-8B as the target, Qwen3-0.6B as the normal S²D AR drafter, and z-lab/Qwen3-8B-DFlash-b16 as the target-conditioned DFlash drafter. The DFlash checkpoint has block size 16, so the DFlash seed always produces 15 speculative tokens. In the split-lookahead hybrid, the SSD side can consume only a prefix of this seed, allowing us to compare $k_{\text{SSD}} = 15$ with $k_{\text{SSD}} = 7$ while keeping the DFlash checkpoint unchanged.

Metrics. We report end-to-end generated-token throughput, decode throughput, average accepted length per verify step, S²D cache hit rate, DFlash seed latency, draft-step latency, and SSD tree-decode latency. Throughput is the user-visible metric. Acceptance and cache hit rate diagnose draft quality and S²D prediction quality. Seed, draft-step, and tree-decode latencies identify whether the bottleneck is DFlash, the AR tree cache, or cross-process orchestration.

3.2. Reference numbers from SSD and DFlash

Table 1 summarizes the published S²D numbers most relevant to this project. The Qwen3-32B results are especially important: they show that S²D still improves over ordinary SD, but the speedup over AR is smaller than in the Llama-70B setting because the verifier is faster and the draft/communication overheads are less easily amortized.

Table 2 gives the DFlash reference point. DFlash is not merely a generic block-diffusion model; it uses target activations and generates the whole block in one conditioned draft pass. Published SGLang results on a single B200 show large Qwen3-8B speedups, which is why DFlash is a more promising diffusion substrate for S²D than iterative block diffusion.

3.3. Our findings

Traditional block diffusion is too slow for SSD. The direct block-diffusion integration did not produce useful S²D speedups. The reason is mechanical: although a diffusion forward predicts all block positions in parallel, the sampler

Table 2. Published DFlash SGLang throughput for Qwen3-8B on a single B200 with FlashAttention-4 [3].

Task	AR	DFlash	Speedup	Avg. τ
Math500	230	1175	5.1×	8.01
HumanEval	229	955	4.2×	6.50

Table 3. Current DFlash+S²D prototype diagnostics. This run used Qwen3-8B target, Qwen3-0.6B AR drafter, and Qwen3-8B-DFlash-b16; it was a small diagnostic run, not the full 512-prompt evaluation.

Metric	Value
Throughput	11.68 tok/s
Average accepted length	3.30 tokens
S ² D cache hit rate	0.60
DFlash seed latency	34.48 ms
SSD tree-decode latency	79 ms

must run many refinement forwards to get acceptable token quality. In the settings where generations were coherent enough to verify well, the number of diffusion passes was close to the number of useful block tokens. The result was a draft-side bottleneck rather than a hidden draft cost. Our best traditional block-diffusion throughput was only about 20 tokens/s, far below the AR and SSD reference points in Table 1.

DFlash improves the diffusion primitive but not yet the full hybrid. The DFlash path removes the main flaw above: it uses fresh target activations and proposes the block in one draft call. Our current hybrid, however, is not yet competitive end to end. On the Qwen3-8B/Qwen3-0.6B/DFlash-b16 prototype, a 32-token smoke run measured:

These numbers should not be interpreted as evidence that DFlash is slow. Rather, they show that our first hybrid pays too much systems overhead for the small verifier used in the test. The DFlash seed itself is tens of milliseconds, while the SSD tree-decode path is even larger. With an 8B verifier on busy A6000s, target verification is not slow enough for the separate-GPU draft path, tensor transfers, cache bookkeeping, and tree construction to be hidden.

The H100 Modal runs let us test one direct systems optimization: keep the DFlash seed length at 15, but reduce only the SSD tree depth to 7. Table 4 compares matched two-H100 runs with seeded random inputs, four requests, 128 input tokens, and 64 output tokens. The shorter SSD lookahead slightly reduces average tokens accepted per step, but it cuts draft-side tree work enough to improve total throughput.

The gain is modest but informative. Reducing k_{SSD} from 15 to 7 improved end-to-end throughput by about 5.4%, decode throughput by about 10%, and nearly halved measured draft-step time. This suggests that tree-cache construction

Table 4. Two-H100 Modal DFlash+S²D split-lookahead diagnostic. DFlash always emits its native 15-token block; k_{SSD} controls how many seed positions the SSD verifier/tree consumes.

k_{SSD}	Total	Decode	Tok./step	Acc. frac.	Hit	Draft ms
15	31.15	40	3.71	0.18	0.68	57.02
7	32.82	44	3.38	0.34	0.71	29.04

Table 5. V1 (DFlash on misses only) vs V2 (DFlash on hits and misses) on Qwen3-8B with Qwen3-0.6B + Qwen3-8B-DFlash-b16. GSM8K, 8×128, greedy, two H100. τ is mean accepted-suffix length.

Metric	V1	V2
Throughput (tok/s)	77.77	31.33
τ overall	6.79	2.10
τ on cache hit / miss	7.36 / 4.04	1.97 / 4.37
Cache hit rate	0.83	0.95

and async draft bookkeeping are important fixed costs in the hybrid. It also explains why simply increasing the speculative block length is not always beneficial: beyond the point where additional drafted positions are rarely accepted, the larger SSD tree can cost more than it saves.

V1 vs V2 on real GSM8K prompts. Re-running the prototype on real GSM8K chat-templated prompts (8 prompts, 128 output tokens, greedy, $k=15$) on the same two-H100 Modal configuration removes the random-token confound and lets us compare V1 against V2 (Sec. 2.4).

V1’s $\tau=6.79$ matches DFlash’s published $\tau=6.54$ on the same target+dataset (Table 2), confirming the integration is correct; V1’s 77.77 tok/s is $\sim 2.5\times$ over plain AR Qwen3-8B on H100. V2 collapses hit-path acceptance from 7.36 to 1.97 while leaving miss-path acceptance unchanged: DFlash is trained on [anchor, MASK, . . .] blocks, so cached tokens as a non-mask prior are out-of-distribution input. Since 95% of steps hit, the broken hit response halves throughput. The V1 design choice of trusting the cache on hits is therefore the right one; making V2 viable would need DFlash retrained on partial-mask blocks.

Small verifiers are the wrong stress test. S²D benefits when the target verifier is much slower than the draft path. If the verifier is small, the latency gap between verifier and drafter shrinks and fixed costs dominate: inter-GPU tensor transfers, process synchronization, PyTorch dispatch, CUDA graph warmup, cache-key matching, and tree-cache population. This explains why the Qwen3-32B SSD speedup reported by Kumar et al. [5] is lower than the Llama-70B speedup, and why our Qwen3-8B prototype is an especially unfavorable testbed.

Compute limited the representative test. The decisive experiment is DFlash+S²D on a larger verifier, such as Qwen3-32B, Qwen3-Coder-30B-A3B, Qwen3.6-35B-A3B, or Llama-70B, with the DFlash and AR draft paths isolated on a separate GPU. We were not able to run that full experiment within the available GPU credits and shared-server availability. This is a practical limitation, not just an engineering inconvenience: the hypothesis specifically concerns latency hiding at large target-model scale, and small smoke runs cannot validate it.

3.4. Discussion

The negative result is still useful. A generic block-diffusion model is not a drop-in win for S²D because iterative denoising is too expensive. DFlash is the more plausible route because it collapses block drafting to one target-conditioned pass and achieves high acceptance in published serving experiments. Our hybrid design also preserves the strongest part of S²D: cache hits still return precomputed AR branches immediately. The open question is whether DFlash can make misses cheap enough on a verifier large enough that the target pass dominates all other overhead.

4. Conclusion

We investigated whether diffusion drafters are a good fit for Speculative² Decoding. The direct approach—replacing the AR drafter with a traditional masked or block-diffusion model—did not work well. Although block diffusion predicts many positions per forward pass, practical generation still requires many refinement passes. In our system this made the diffusion drafter, not the verifier, the bottleneck, with throughput around 20 tokens/s.

DFlash is a better match to the problem because it uses target activations from the verifier and drafts a full block in one conditioned pass. We implemented a DFlash-seeded S²D hybrid in which DFlash serves speculation-cache misses and the normal AR S²D drafter serves cache hits and prepares the future tree cache. This preserves exact verification and keeps the fast S²D hit path intact. Our initial local Qwen3-8B prototype was not fast, reaching 11.68 tokens/s with 34.48 ms DFlash seed latency and 79 ms SSD tree-decode latency, but it isolated the right next bottlenecks. Moving the same prototype to a Modal H100 environment and decoupling the DFlash seed length from the SSD lookahead improved a matched two-H100 diagnostic from 31.15 to 32.82 tokens/s, mostly by reducing draft-side tree work. This is not a final serving result, but it makes the bottleneck more specific: the hybrid is dominated by tree-cache construction and orchestration unless the verifier is large enough to hide those costs.

On real GSM8K prompts the same two-H100 prototype reached 77.77 tok/s with $\tau=6.79$, matching DFlash’s published acceptance for Qwen3-8B (Table 5). A universal-

drafter ablation that also runs DFlash on cache hits halved throughput, because non-mask priors are out-of-distribution for DFlash. DFlash adds value precisely where the cache has no answer; trusting the cache on hits is the right design.

The broader conclusion is that S^2D is sensitive to the verifier/drafter latency ratio. On small target models, cross-GPU communication and bookkeeping overhead can erase any benefit from asynchronous drafting. Future work should run the DFlash hybrid on a larger verifier—Qwen3-32B, Qwen3-Coder-30B-A3B, Qwen3.6-35B-A3B, or Llama-70B—using enough GPU memory to keep the verifier and draft paths isolated. That experiment requires more compute than was available for this project, but it is the appropriate test of whether DFlash can make S^2D cache misses cheap enough to improve the full throughput-latency frontier.

References

- [1] Marianne Arriola, Aaron Gokaslan, Justin T Chiu, Zhihan Yang, Zhixuan Qi, Jiaqi Han, Subham Sekhar Sahoo, and Volodymyr Kuleshov. Block diffusion: Interpolating between autoregressive and diffusion language models. In *ICLR*, 2025. [1](#), [2](#)
- [2] Charlie Chen, Sebastian Borgeaud, Geoffrey Irving, Jean-Baptiste Lespiau, Laurent Sifre, and John Jumper. Accelerating large language model decoding with speculative sampling. *arXiv preprint arXiv:2302.01318*, 2023. [1](#), [2](#)
- [3] Jian Chen, Yesheng Liang, and Zhijian Liu. DFlash: Block diffusion for flash speculative decoding, 2026. [2](#), [4](#)
- [4] Marjan Ghazvininejad, Omer Levy, Yinhan Liu, and Luke Zettlemoyer. Mask-predict: Parallel decoding of conditional masked language models. In *Empirical Methods in Natural Language Processing (EMNLP)*, 2019. [2](#)
- [5] Tanishq Kumar, Tri Dao, and Avner May. Speculative speculative decoding, 2026. [1](#), [2](#), [3](#), [4](#), [5](#)
- [6] Yaniv Leviathan, Matan Kalman, and Yossi Matias. Fast inference from transformers via speculative decoding. In *International Conference on Machine Learning (ICML)*, 2023. [1](#), [2](#)
- [7] Subham Sekhar Sahoo, Marianne Arriola, Yair Schiff, Aaron Gokaslan, Edgar Marroquin, Justin T Chiu, Alexander Rush, and Aditya Grover. Simple and effective masked diffusion language models. In *NeurIPS*, 2024. [1](#), [2](#)
- [8] An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, Chujie Zheng, et al. Qwen3 technical report, 2025. [3](#)
- [9] Lianmin Zheng, Liangsheng Yin, Zhiqiang Xie, Chuyue Livia Sun, Jeff Huang, Cody Hao Yu, Shiyi Cao, Christos Kozyrakis, Ion Stoica, Joseph E. Gonzalez, Clark Barrett, and Ying Sheng. SGLang: Efficient execution of structured language model programs. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. [2](#)